

In [1]:

```
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K
```

Using TensorFlow backend.

In [2]:

```
K.image_data_format()
```

Out[2]:

```
'channels_last'
```

In [3]:

```
(x_train, y_train), (x_test, y_test) = mnist.load_data()
x_train = x_train / 255.0
x_test = x_test / 255.0
```

In [4]:

```
print(x_train.shape)
```

```
(60000, 28, 28)
```

In [5]:

```
print(x_test.shape)
```

```
(10000, 28, 28)
```

## Reshape for CNN

In [6]:

```
x_train = x_train.reshape(-1, 28, 28, 1)    #Reshape for CNN - should work!!
print(x_train.shape)
```

```
(60000, 28, 28, 1)
```

In [7]:

```
x_test = x_test.reshape(-1, 28, 28, 1)
print(x_test.shape)
```

```
(10000, 28, 28, 1)
```

In [8]:

```
input_shape1 = (28, 28, 1)
```

## Hyper paramemters

In [9]:

```
nb_epoch = 5
num_classes = 10
batch_size = 128
```

## Model definition

In [10]:

```
model = Sequential()
model.add(Conv2D(32, kernel_size=(3, 3),
                 activation='relu',
                 input_shape=input_shape))
model.add(Conv2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))
model.add(Dropout(0.25))
model.add(Flatten())
model.add(Dense(128, activation='relu'))
model.add(Dropout(0.5))
model.add(Dense(num_classes, activation='softmax'))
```

In [11]:

```
model.summary()
```

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 26, 26, 32)	320
conv2d_2 (Conv2D)	(None, 24, 24, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(None, 12, 12, 64)	0
dropout_1 (Dropout)	(None, 12, 12, 64)	0
flatten_1 (Flatten)	(None, 9216)	0
dense_1 (Dense)	(None, 128)	1179776
dropout_2 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 10)	1290

=====  
Total params: 1,199,882  
Trainable params: 1,199,882  
Non-trainable params: 0  
=====

Q. why 1,199,882 ?

In [12]:

```
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
```

In [13]:

```
model.fit(x_train, y_train, epochs=nb_epoch)
```

```
Epoch 1/5
60000/60000 [=====] - 125s 2ms/step - loss:
0.1902 - acc: 0.9421
Epoch 2/5
60000/60000 [=====] - 121s 2ms/step - loss:
0.0786 - acc: 0.9765
Epoch 3/5
60000/60000 [=====] - 124s 2ms/step - loss:
0.0601 - acc: 0.9823
Epoch 4/5
60000/60000 [=====] - 122s 2ms/step - loss:
0.0492 - acc: 0.9849
Epoch 5/5
60000/60000 [=====] - 122s 2ms/step - loss:
0.0415 - acc: 0.9869
```

## ***Real World Challenge: Difference between training and testing set accuracy***

### **Test accuracy**

In [17]:

```
score = model.evaluate(x_test, y_test)
```

```
10000/10000 [=====] - 7s 667us/step
```

In [19]:

```
score
```

Out[19]:

```
[0.032295464331466066, 0.9896]
```

In [20]:

```
print('Test loss:{}'.format(score[0]))
print('Test accuracy:{}'.format(score[1]))
```

```
Test loss:0.032295464331466066
Test accuracy:0.9896
```

### **Training accuracy**

In [21]:

```
score = model.evaluate(x_train, y_train)
```

```
60000/60000 [=====] - 38s 640us/step
```

In [22]:

```
print('Training loss:{}'.format(score[0]))  
print('Training accuracy:{}'.format(score[1]))
```

Training loss:0.011586855655677695

Training accuracy:0.9960166666666667

**Q. What is the difference between the training and test accuracy ?**